

Chapter 6: Architecture

Floating-Point Instructions

RISC-V Floating-Point Extensions

- RISC-V offers three floating point extensions:
 - **RVF:** single-precision (32-bit)
 - 8 exponent bits, 23 fraction bits
 - **RVD:** double-precision (64-bit)
 - 11 exponent bits, 52 fraction bits
 - **RVQ:** quad-precision (128-bit)
 - 15 exponent bits, 112 fraction bits

Floating-Point Registers

- 32 Floating point registers
- **Width** is highest precision – for example, if RVQ is implemented, registers are 128 bits wide
- When multiple floating point extensions are implemented, the lower-precision values occupy the lower bits of the register

Floating-Point Registers

Name	Register Number	Usage
ft0-7	f0-7	Temporary variables
fs0-1	f8-9	Saved variables
fa0-1	f10-11	Function arguments/Return values
fa2-7	f12-17	Function arguments
fs2-11	f18-27	Saved variables
ft8-11	f28-31	Temporary variables

Floating-Point Instructions

- Append `.s` (single), `.d` (double), `.q` (quad) for precision. I.e., `fadd.s`, `fadd.d`, and `fadd.q`
- **Arithmetic operations:**
`fadd`, `fsub`, `fdiv`, `fsqrt`, `fmin`, `fmax`, multiply-add (`fmadd`, `fmsub`, `fnmadd`, `fnmsub`)
- **Other instructions:**
`move` (`fmv.x.w`, `fmv.w.x`)
`convert` (`fcvt.w.s`, `fcvt.s.w`, etc.)
`comparison` (`feq`, `flt`, `fle`)
`classify` (`fclass`)
`sign injection` (`fsgnj`, `fsgnjn`, `fsgnjx`)

See Appendix B for additional RISC-V floating-point instructions.

Floating-Point Multiply-Add

- `fmaddd` is the most critical instruction for signal processing programs.
- Requires four registers.

```
fmaddd.f f1, f2, f3, f4    # f1 = f2 x f3 + f4
```

Floating-Point Example

C Code

```
int i;  
float scores[200];  
  
for (i=0; i<200; i=i+1)  
  
    scores[i]=scores[i]+10;
```


RISC-V assembly code

```
# s0 = scores base address, s1 = i  
addi s1, zero, 0          # i = 0  
addi t2, zero, 200        # t2 = 200  
addi t0, zero, 10         # ft0 = 10.0  
fcvt.s.w ft0, t0  
  
for:  
    bge s1, t2, done      # i>=200? done  
    slli t0, s1, 2        # t0 = i*4  
    add t0, t0, s0        # scores[i] address  
    flw ft1, 0(t0)        # ft1=scores[i]  
    fadd.s ft1, ft1, ft0   # ft1=scores[i]+10  
    fsw ft1, 0(t0)        # scores[i] = t1  
    addi s1, s1, 1        # i = i+1  
    j for                 # repeat  
done:
```

Floating-Point Instruction Formats

- Use R-, I-, and S-type formats
- Introduce another format for multiply-add instructions that have 4 register operands: R4-type

R4-Type



Chapter 6: Architecture

Exceptions

Exceptions

- Unscheduled function call to *exception handler*
- Caused by:
 - Hardware, also called an *interrupt*, e.g., keyboard
 - Software, also called *traps*, e.g., undefined instruction
- When exception occurs, the processor:
 - Records the cause of the exception
 - Jumps to exception handler
 - Returns to the program

Exception Causes

Exception	Cause
Instruction address misaligned	0
Instruction access fault	1
Illegal instruction	2
Breakpoint	3
Load address misaligned	4
Load access fault	5
Store address misaligned	6
Store access fault	7
Environment call from U-Mode	8
Environment call from S-Mode	9
Environment call from M-Mode	11

RISC-V Privilege Levels

- In RISC-V, exceptions occur at various **privilege levels**.
- Privilege levels limit access to memory or certain (privileged) instructions.
- RISC-V privilege modes are (from highest to lowest):
 - **Machine** mode (bare metal)
 - **System** mode (operating system)
 - **User** mode (user program)
 - **Hypervisor** mode (to support virtual machines)
- For example, a program running in M-mode (machine mode) can access all memory or instructions – it has the highest privilege level.

Exception Registers

- Each privilege level has registers to handle exceptions
- These registers are called control and status registers (**CSRRs**)
- We discuss **M-mode** (machine mode) exceptions, but other modes are similar
- M-mode registers used to handle exceptions are:
 - `mtvec`, `mcause`, `mepc`, `mscratch`

(Likewise, S-mode exception registers are: `stvec`, `scause`, `sepc`, and `mscratch`; and so on for the other modes.)

Exception Registers

- CSRRs are not part of register file
- M-mode CSRRs used to handle exceptions
 - **mtvec**: holds address of exception handler code
 - **mcause**: Records cause of exception
 - **mepc** (Exception PC): Records PC where exception occurred
 - **mscratch**: scratch space in memory for exception handlers

Exception-Related Instructions

- Called **privileged instructions** (because they access CSRRs)
 - **csrr**: CSR register read
 - **csrw**: CSR register write
 - **csrrw**: CSR register read/write
 - **mret**: returns to address held in mepc
- **Examples:**

```
csrr t1, mcause           # t1 = mcause
csrw mepc, t2             # mepc = t2
csrrw t0, mscratch, t1    # t0 = mscratch
                          # mscratch = t1
```

Exception Handler Summary

- When a processor **detects an exception**:
 - It **jumps to exception handler** address in `mtvec`
 - The exception handler then:
 - **saves registers** on small stack pointed to by `mscratch`
 - Uses `csrr` (CSR read) to **look at cause** of exception (in `mcause`)
 - **Handles exception**
 - When finished, optionally **increments mepc by 4** and **restores registers** from memory
 - And then either **aborts** the program or **returns to user code** (using `mret`, which returns to address held in `mepc`)

Example Exception Handler Code

- Check for **two types of exceptions**:
 - Illegal instruction (`mcause = 2`)
 - Load address misaligned (`mcause = 4`)

Example Exception Handler Code

```
# save registers that will be overwritten
csrrw t0, mscratch, t0    # swap t0 and mscratch
sw     t1, 0(t0)          # [mscratch] = t1
sw     t2, 4(t0)          # [mscratch+4] = t2

# check cause of exception
csrr   t1, mcause         # t1=mcause
addi   t2, x0, 2          # t2=2 (illegal instruction exception code)
illegalinstr:
bne    t1, t2, checkother # branch if not an illegal instruction
csrr   t2, mepc           # t2=exception PC
addi   t2, t2, 4          # increment exception PC
csrw   mepc, t2           # mepc=t2
j      done              # restore registers and return
checkother:
addi   t2, x0, 4          # t2=4 (load address misaligned exception code)
bne    t1, t2, done       # branch if not a misaligned load
j      exit              # exit program

# restore registers and return from the exception
done:
lw     t1, 0(t0)          # t1 = [mscratch]
lw     t2, 4(t0)          # t2 = [mscratch+4]
csrrw t0, mscratch, t0    # swap t0 and mscratch
mret                                # return to program
exit:
...
```

Checks for two types of exceptions:

- **Illegal instruction**
(mcause = 2)
- **Load address misaligned**
(mcause = 4)